

Toxic Comment Classification Challenge

Problem Statement - You are provided with a large number of Wikipedia comments which have been labeled by human raters for toxic behavior. The types of toxicity are toxic, severe_toxic, obscene, threat, insult and identity_hate. You must create a model which predicts a probability of each type of toxicity for each comment.

Dataset Description - We are given with a dataset having the following dataset description -

- ID's of comments
- Text Comments
- Six output columns corresponding to each output class having binary class as '0' & '1'.
- The output columns are as follows:
 - - toxic
 - - severe_toxic
 - - obscene
 - - threat
 - - insult
 - - identity_hate
- The dataset has 159571 comment text data and corresponding output values.

Understanding the Problem Statement - We are given with the commented text and the corresponding labels and using supervised machine learning models on the training data we should be able to correctly predict the probabilities of the given comment lying in a particular output class.

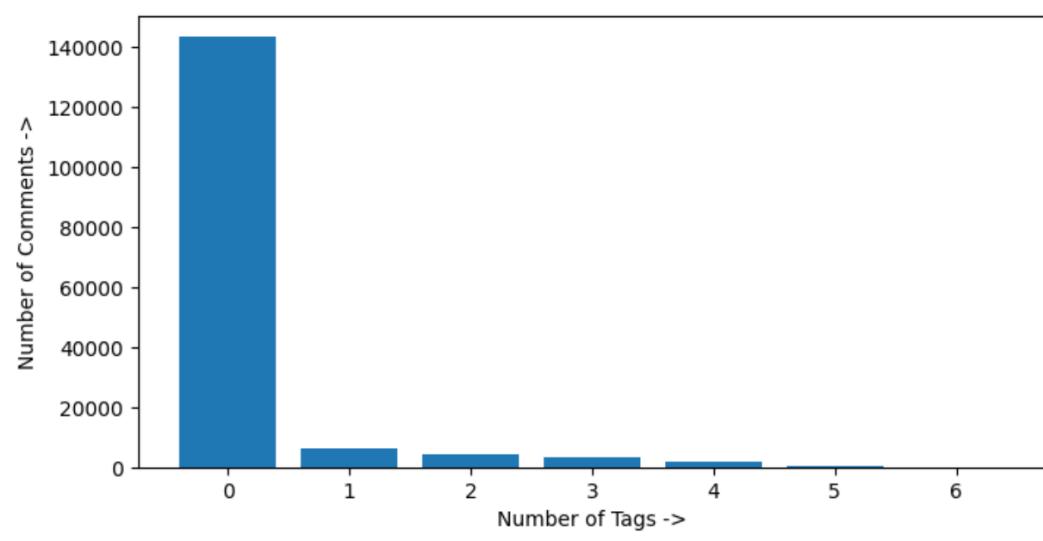
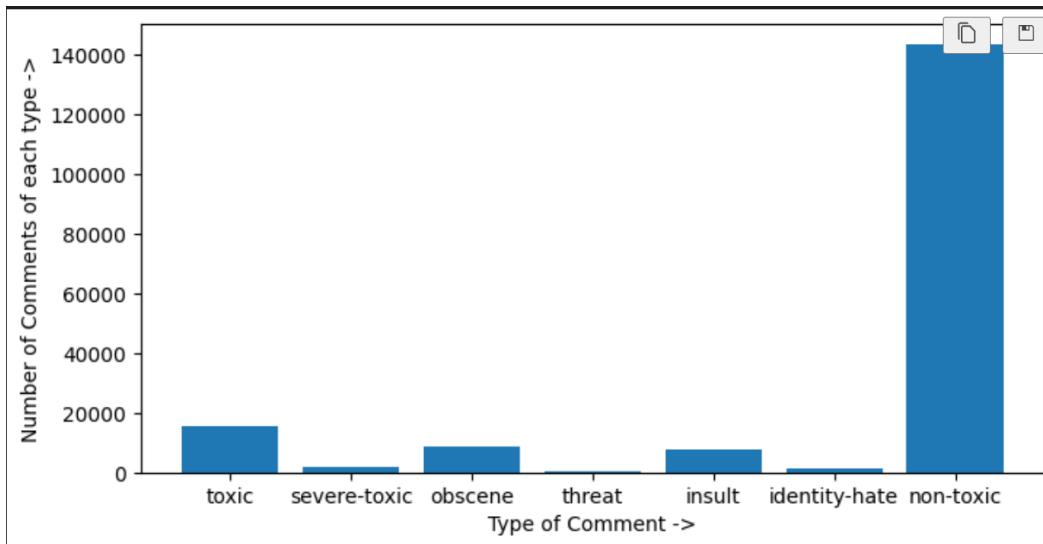
Proposed solution flow - On observing the dataset, we observe that we need to deal with the text data and to train the various models on it. To work further we can't directly continue with the given text data. So the workflow pipeline will be as follows:-

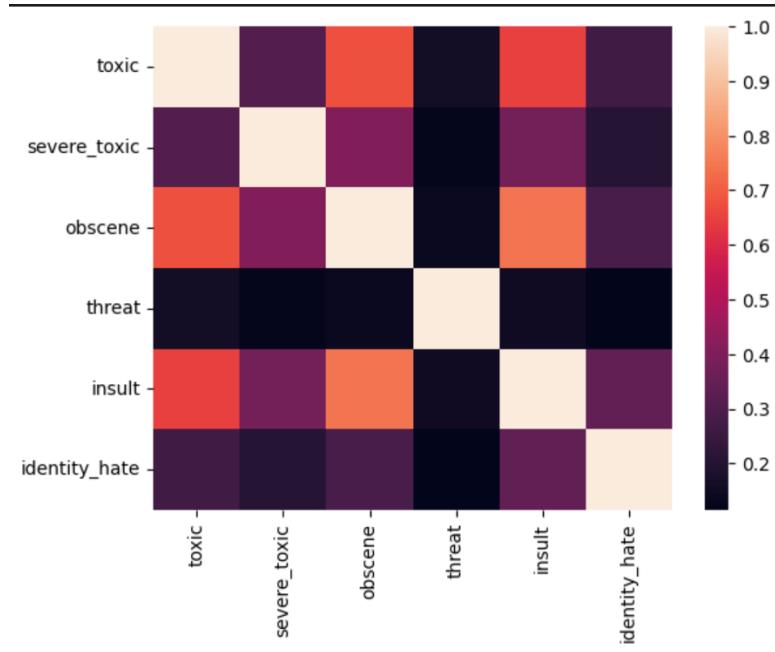
- Preprocess the text data.
- Balance out the whole data if required to remove biasness
- Train various ML models and get the evaluation reports
- Find the best suitable ML models and the best method to get the final predictions.
- Predict the probabilities of each comment to be in any particular class for test data.
- Create the final pipeline on the basis of a trained model so that any random comment can be classified.

Visualizing and Analyzing the data

- We got the dataframe and checked for the **NULL values** in the complete dataframe and got the output that there are no NULL values in the given data.
- To analyze the data completely and get a complete idea about the dataframe, we obtained the counts for each output class of the number of '1' values for a particular column which is the presence of that particular class output in the whole data.

```
Total comments = 159571
Total non_toxic comments = 143346
Count of each type of comment:
toxic           15294
severe_toxic    1595
obscene          8449
threat            478
insult            7877
identity_hate    1405
```





- We obtained the **confusion matrix** considering any two columns and obtained all the confusion matrices.

Analyze the comments

- From the given dataframe, we obtained the particular **column comments** belonging to each particular class, i.e class output is '1'.

Visualize the comments

- Using “World-Cloud”, a **visualization** method that displays how frequently words appear in a given body of text, by making the size of each word proportional to its frequency.

Frequent words in Non-toxic Comments



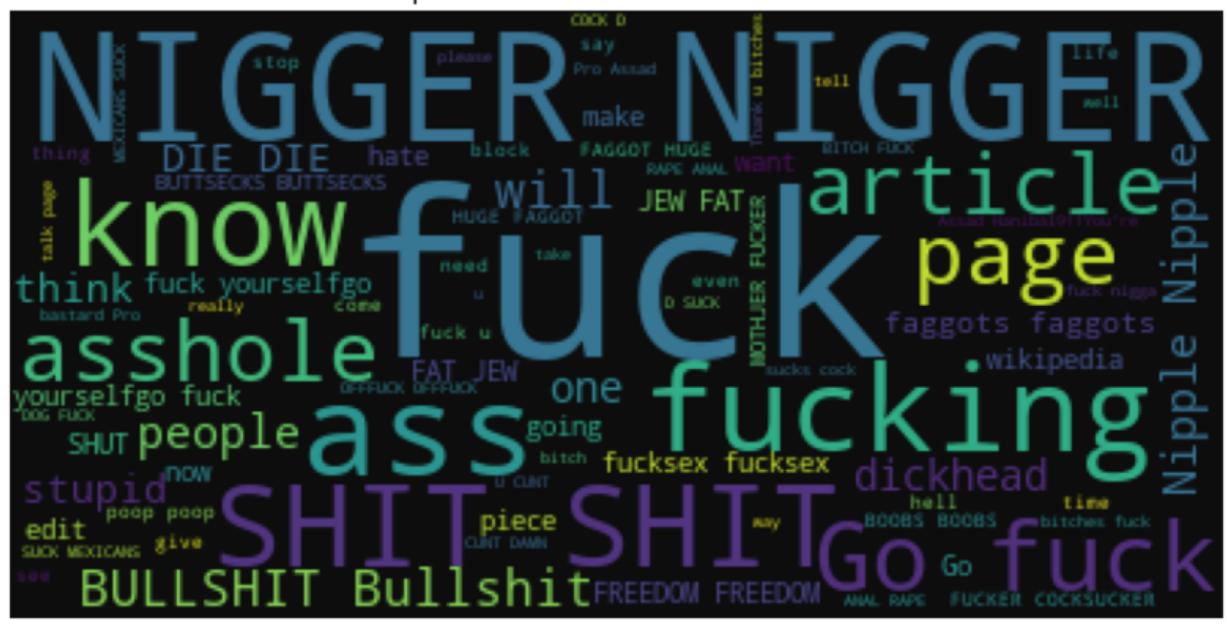
Frequent words in Toxic Comments



Frequent words in Severe-toxic Comments



Frequent words in Obscene Comments



Frequent words in Threat Comments



Frequent words in Insult Comments



Frequent words in Identity Hate Comments



Preprocessing

- **Cleaning Text** - First we have to clean our text comments so that our tokenization and vectorizer will work better, we have to remove the extra white spaces, punctuation signs, numbers, email addresses and all the unnecessary things in the text comments which is not affecting our classification.
 - **Apply Tokenizer** - Tokenization is the process of breaking down a text document or a sentence into individual words or tokens. In other words, tokenization is a technique used to segment text into meaningful units called tokens. The tokens can be words, numbers, or punctuation marks. So we are applying the tokenizer on the text comments to get a tokenized set and further proceed with different cleaning techniques.

- **Removing stop words** - Stop words like "the", "a", "an", "and", "or", "of", "in" are removed from the tokenized data obtained using “nltk”, i.e. words in corpus of nltk.

Balancing out the data

We are generating a balanced dataset by augmenting the minority class in the original dataset. It follows the following steps:

1. Splits the original dataset into a training and testing set, with 80% of the data used for training and 20% for testing.
2. For each output column (toxic, severe_toxic, obscene, threat, insult, identity_hate), the code identifies the minority class and the number of samples in that class.
3. It calculates the number of samples needed to be added to balance the dataset, which is the difference between the number of samples in the minority and majority class divided by the number of output columns (to ensure a balanced dataset across all output columns).
4. The code generates augmented data for the minority class by randomly selecting samples from the majority class and flipping the output value for the minority class.
5. The augmented data is appended to the training set.
6. If the size of the training set exceeds the target size, the code randomly samples a subset of the data to reduce it to the target size.
7. The balanced dataset is saved to a new file and returned as a pandas dataframe.
8. Note that this code assumes that the input dataset is already preprocessed (e.g., tokenized, stop words removed, and vectorized). It only focuses on balancing the dataset by augmenting the minority class. And thus we are providing the dataframe to the code after dropping the other useful columns for re-sampling.

Vectorising the text data

Vectorizing text data refers to the process of converting raw text data into numerical form so that it can be processed by machine learning algorithms. We have used “TfidfVectorizer” which takes in a corpus of documents and generates a matrix of TF-IDF features, where each row represents a document and each column represents a word in the vocabulary. The vectorizer will give us the numerical features out of our text data, so that we can train our models on these numerical features.

Train-Test Split for final evaluation of the model

We now have obtained the two dataframes finally. The first one is the original dataframe and the second one is the data frame created by balancing and reducing the original dataframe.

Codeflow afterwards

Now we are applying different machine learning models on the obtained data frames and further on the basis of the evaluation reports provided by standard sklearn library, we will make the final decision for which we will go with the final dataframe and models for getting final predictions.

Function for models training and testing and saving:

Here we have used many models for classification for each output column, so for each output column we have called the particular model function to train the particular model on each output column using a loop the function to call the respective function and to train the model on each column. We

used the “pickle.dump” function to save the trained model by giving a unique string address which is generated by our function. We have also used the class weights parameter for increasing the recall of our model. This function will also give a classification report including precision, recall, accuracy on the test data which we are taking by using test train split on the original train.csv file.

And we have taken the hyperparameters of some models from a research paper*.

We made the this function for each model which includes:

- Random Forest
- Logistic Regression
- Support Vector Machine (SVM)
- K Nearest Neighbors (KNN)
- Gradient Boosting Classifier (GBM)
- Quadratic Discriminant Analysis (QDA)
- XGBoost Classifier
- Light GBM Classifier

Making the Models:

We have obtained the 2D array of model names and short codes which individual array will contain the names of models which are giving the best results for that particular output column. So by using the same address of the models which we saved earlier we called our saved models using the pickle.load() function.

The models for each output column that we have selected by analyzing the precision, recall, accuracy and f1-score are:

Severe -> Random Forest, Logistic Regression, Support Vector Machine (SVM), K Nearest Neighbors (KNN), XGBoost Classifier, Light GBM Classifier.

Severe_toxic -> Random Forest, Logistic Regression, Support Vector Machine (SVM), XGBoost Classifier, Light GBM Classifier.

Obscene -> Logistic Regression, Support Vector Machine (SVM), XGBoost Classifier, Light GBM Classifier.

Threat -> Logistic Regression, Support Vector Machine (SVM), XGBoost Classifier, Light GBM Classifier.

Insult -> Random Forest, Logistic Regression, Support Vector Machine (SVM), Gradient Boosting Classifier (GBM), XGBoost Classifier, Light GBM Classifier.

Identity_hate -> Random Forest, Support Vector Machine (SVM), XGBoost Classifier.

Function to predict the probability of each output column

This function basically takes the **string input** and provides the **probability** of that particular comment text to be in a particular **targeted output column**.

We made a function pred_prob which takes a string input like as a comment on which we have to find the probability of each output column, and it is taking an output array obtained from the particular model for that particular output column which contains all the models which are the best. And then we have used the **voting** for all the outputs obtained for all the best models to find the probability of the particular column output.

For the prediction of the models, we have followed the standard preprocessing with text & get the predictions using the final trained models.

This function will give a single probability value of that particular column label, then we call this function for each column by passing the particular best models of this column and appending the probability to an array. Then find the output binary classes by using the probability we got, if probability is greater than equal to 0.5, then we will give the output for that particular output class '1' otherwise 0.

Judgment on dataframes and its corresponding models

Now we have multiple dataframes and their corresponding trained models available. Now on the basis of the accuracy, precision, recall & f1 scores available for the models, we concluded to go further with the original data because the balanced data was synthetically created and may or not able to get the desired outputs every time.

Following are the evaluation results for various models on the augmented dataset :-

1) Random forest					
94.13097936299654					
	precision	recall	f1-score	support	
0	0.95	0.98	0.97	7293	
1	0.84	0.68	0.75	1090	
94.13097936299654					
	precision	recall	f1-score	support	
0	0.96	0.97	0.97	7338	
1	0.78	0.73	0.76	1045	
91.36347369676727					
	precision	recall	f1-score	support	
0	0.94	0.96	0.95	7076	
1	0.75	0.67	0.71	1307	
91.32768698556603					
	precision	recall	f1-score	support	
0	0.95	0.95	0.95	7341	
1	0.66	0.64	0.65	1042	
87.39114875342956					
	precision	recall	f1-score	support	
0	0.93	0.92	0.93	7070	
1	0.60	0.60	0.60	1313	
82.6434450673983					
	precision	recall	f1-score	support	
0	0.93	0.86	0.90	7297	
1	0.39	0.58	0.47	1086	

1) Random forest					
94.13097936299654					
	precision	recall	f1-score	support	
0	0.95	0.98	0.97	7293	
1	0.84	0.68	0.75	1090	
94.13097936299654					
	precision	recall	f1-score	support	
0	0.96	0.97	0.97	7338	
1	0.78	0.73	0.76	1045	
91.36347369676727					
	precision	recall	f1-score	support	
0	0.94	0.96	0.95	7076	
1	0.75	0.67	0.71	1307	
91.32768698556603					
	precision	recall	f1-score	support	
0	0.95	0.95	0.95	7341	
1	0.66	0.64	0.65	1042	
87.39114875342956					
	precision	recall	f1-score	support	
0	0.93	0.92	0.93	7070	
1	0.60	0.60	0.60	1313	
82.6434450673983					
	precision	recall	f1-score	support	
0	0.93	0.86	0.90	7297	
1	0.39	0.58	0.47	1086	

2) logistic regression					
92.30455898480339					
	precision	recall	f1-score	support	
0	0.98	0.93	0.96	28882	
1	0.56	0.84	0.68	3033	
96.63481121729595					
	precision	recall	f1-score	support	
0	1.00	0.97	0.98	31591	
1	0.22	0.87	0.35	324	
96.07081309728967					
	precision	recall	f1-score	support	
0	0.99	0.97	0.98	30215	
1	0.59	0.88	0.70	1700	
98.74980416731944					
	precision	recall	f1-score	support	
0	1.00	0.99	0.99	31826	
1	0.15	0.72	0.24	89	
94.37568541438195					
	precision	recall	f1-score	support	
0	0.99	0.95	0.97	30346	
1	0.46	0.87	0.60	1569	
95.56948143506189					
	precision	recall	f1-score	support	
0	1.00	0.96	0.98	31619	
1	0.15	0.79	0.25	296	

3) SVM					
81.59370153882858					
	precision	recall	f1-score	support	
0	0.96	0.82	0.89	7293	
1	0.39	0.76	0.52	1090	
70.4759632589765					
	precision	recall	f1-score	support	
0	0.94	0.71	0.81	7338	
1	0.25	0.69	0.37	1045	
71.16783967553381					
	precision	recall	f1-score	support	
0	0.92	0.72	0.81	7076	
1	0.30	0.65	0.41	1307	
65.83561970654897					
	precision	recall	f1-score	support	
0	0.93	0.66	0.77	7341	
1	0.21	0.64	0.32	1042	
67.57723965167601					
	precision	recall	f1-score	support	
0	0.90	0.69	0.78	7070	
1	0.27	0.61	0.37	1313	
62.089943934152444					
	precision	recall	f1-score	support	
0	0.91	0.63	0.74	7297	
1	0.19	0.57	0.28	1086	

4) knn					
92.77455741814194					
		precision	recall	f1-score	support
0	0.93	0.99	0.96	28882	
1	0.79	0.33	0.46	3033	

98.92527024909917					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.99	1.00	0.99	31591	
1	0.42	0.16	0.23	324	

96.08021306595644					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.96	0.99	0.98	30215	
1	0.80	0.35	0.49	1700	

99.70233432555224					
precision recall f1-score support					
		precision	recall	f1-score	support
0	1.00	1.00	1.00	31826	
1	0.38	0.10	0.16	89	

95.97681341062196					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.97	0.99	0.98	30346	
1	0.70	0.32	0.44	1569	

99.12266959110137					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.99	1.00	1.00	31619	
1	0.64	0.12	0.20	296	

4) KNN					
90.87438864368364					
		precision	recall	f1-score	support
0	0.92	0.98	0.95	7293	
1	0.75	0.45	0.56	1090	

92.92616008588811					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.96	0.96	0.96	7338	
1	0.73	0.69	0.71	1045	

89.01347966121914					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.93	0.94	0.94	7076	
1	0.66	0.62	0.64	1307	

89.97972086365263					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.93	0.96	0.94	7341	
1	0.63	0.46	0.53	1042	

85.24394608135512					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.89	0.94	0.91	7070	
1	0.54	0.39	0.45	1313	

84.65942979840153					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.90	0.92	0.91	7297	
1	0.39	0.33	0.36	1086	

5) XGB classifier					
95.40968196772678					
		precision	recall	f1-score	support
0	0.96	0.99	0.98	28859	
1	0.90	0.59	0.71	3056	

98.98793670687765					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.99	1.00	0.99	31608	
1	0.44	0.20	0.27	307	

97.84114052953157					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.98	0.99	0.99	30204	
1	0.88	0.69	0.78	1711	

99.74620084599718					
precision recall f1-score support					
		precision	recall	f1-score	support
0	1.00	1.00	1.00	31828	
1	0.59	0.23	0.33	87	

97.11107629641235					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.98	0.99	0.98	30320	
1	0.81	0.55	0.66	1595	

99.23546921510261					
precision recall f1-score support					
		precision	recall	f1-score	support
0	0.99	1.00	1.00	31650	
1	0.62	0.21	0.31	265	

7) XGB					
91.86448765358463					
		precision	recall	f1-score	support
0	0.92	0.99	0.96	7349	
1	0.87	0.40	0.55	1034	

89.01347966121914					

6) LGBM					
93.56415478615071					
	precision	recall	f1-score	support	
0	0.98	0.95	0.96	28860	
1	0.63	0.80	0.70	3055	

96.36221212595957					
	precision	recall	f1-score	support	
0	1.00	0.96	0.98	31612	
1	0.19	0.88	0.31	303	

96.550211499295					
	precision	recall	f1-score	support	
0	0.99	0.97	0.98	30217	
1	0.63	0.87	0.73	1698	

98.51167162776125					
	precision	recall	f1-score	support	
0	1.00	0.99	0.99	31815	
1	0.15	0.77	0.24	100	

94.90835030549898					
	precision	recall	f1-score	support	
0	0.99	0.95	0.97	30324	
1	0.49	0.84	0.62	1591	

95.5851480495065					
	precision	recall	f1-score	support	
0	1.00	0.96	0.98	31644	
1	0.14	0.80	0.24	271	

8) LGBM					
89.7411427889777					
	precision	recall	f1-score	support	
0	0.97	0.92	0.94	7349	
1	0.56	0.77	0.65	1034	

87.68937134677323					
	precision	recall	f1-score	support	
0	0.97	0.89	0.93	7336	
1	0.50	0.82	0.62	1047	

85.98353811284743					
	precision	recall	f1-score	support	
0	0.96	0.87	0.91	7036	
1	0.54	0.79	0.64	1347	

84.46856733866159					
	precision	recall	f1-score	support	
0	0.96	0.86	0.91	7358	
1	0.42	0.76	0.54	1025	

80.67517595133008					
	precision	recall	f1-score	support	
0	0.94	0.82	0.88	7038	
1	0.44	0.73	0.55	1345	

77.5020875581534					
	precision	recall	f1-score	support	
0	0.95	0.79	0.86	7340	
1	0.31	0.68	0.43	1043	

7) Decision tree					
93.56415478615071					
	precision	recall	f1-score	support	
0	0.98	0.95	0.96	28860	
1	0.63	0.80	0.70	3055	

96.36221212595957					
	precision	recall	f1-score	support	
0	1.00	0.96	0.98	31612	
1	0.19	0.88	0.31	303	

96.550211499295					
	precision	recall	f1-score	support	
0	0.99	0.97	0.98	30217	
1	0.63	0.87	0.73	1698	

98.51167162776125					
	precision	recall	f1-score	support	
0	1.00	0.99	0.99	31815	
1	0.15	0.77	0.24	100	

94.90835030549898					
	precision	recall	f1-score	support	
0	0.99	0.95	0.97	30324	
1	0.49	0.84	0.62	1591	

95.5851480495065					
	precision	recall	f1-score	support	
0	1.00	0.96	0.98	31644	
1	0.14	0.80	0.24	271	

6) Decision Tree Classifier					
90.14672551592508					
	precision	recall	f1-score	support	
0	0.94	0.94	0.94	7312	
1	0.62	0.61	0.61	1071	

86.85434808541095					
	precision	recall	f1-score	support	
0	0.92	0.93	0.92	7305	
1	0.49	0.47	0.48	1078	

86.55612549206728					
	precision	recall	f1-score	support	
0	0.92	0.92	0.92	7054	
1	0.58	0.57	0.57	1329	

85.81653346057497					
	precision	recall	f1-score	support	
0	0.93	0.91	0.92	7391	
1	0.41	0.45	0.43	992	

83.41882381009185					
	precision	recall	f1-score	support	
0	0.89	0.91	0.90	7077	
1	0.46	0.41	0.44	1306	

44.733389001550755					
	precision	recall	f1-score	support	
0	0.95	0.39	0.55	7309	
1	0.17	0.86	0.28	1074	

Final Output

Then we called the predict_prob function for the initial 1000 of the test.csv file to find the corresponding labels and probabilities of the test.csv comments because the test.csv file has around 1.5 lakh comments and it is taking very much time to predict all the 1.5 lakh comments, and then we made the final dataframe to show the output which includes our id column, comment_text column and the output label probabilities and classes. And then at the end we save our dataframe to the google drive to download my result.csv file to upload that.

Link for our model results →

https://drive.google.com/drive/folders/1VjPzA_F612-AT6GVxXp2czKCWjX_0y6k?usp=sharing

